



# <epam> Docker for java developers Day 1



- Dmitry Buhtiyarov, Siarhei Beliakou, 2018



# Agenda

## Containers Overview:

- Docker Components
- Docker Architecture
- Docker Storage Drivers


## Installation and Configuration

## Creating Docker Images

## Running Containers

## Getting Logs From Containers

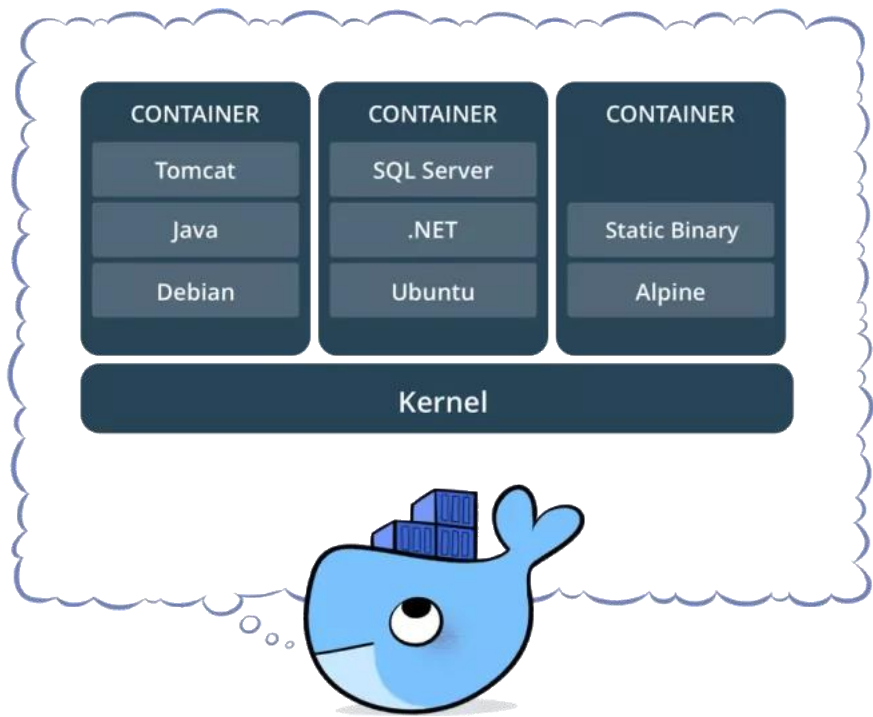
## Demo



## CONTAINERS OVERVIEW

- **What Containers and Images are about**
- **Virtual Machines vs Containers**
- **Containers History**
- **Containers Underlying Technologies**
- **Docker Architecture**
- **Storage Drivers**

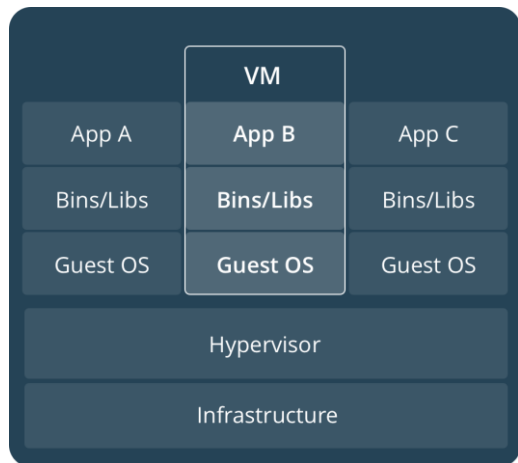
# Virtual Machines vs Containers



A **container** is something quite similar to a virtual machine, which can be used to **contain and execute** all the software required to run a particular program or set of programs.

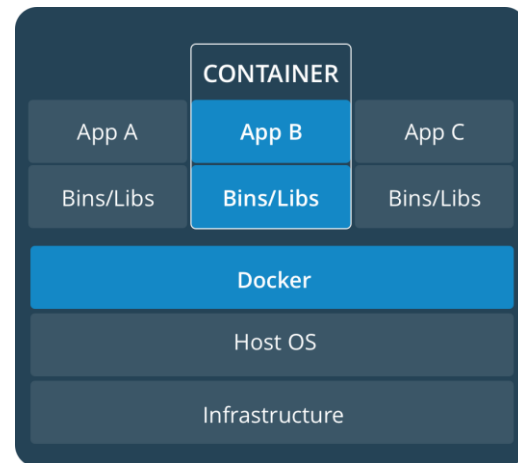
The container includes an operating system (typically some flavor of Linux) as base, plus any software installed on top of the OS that might be needed. This container can therefore be run as a self-contained virtual environment, which makes it a lot easier to reproduce the same analysis on any infrastructure that supports running the container, from your laptop to a cloud platform, without having to go through the pain of identifying and installing all the software dependencies involved. You can even have multiple containers running on the same machine, so you can easily switch between different environments if you need to run programs that have incompatible system requirements.

# Virtual Machines vs Containers



## VIRTUAL MACHINES

Virtual machines (VMs) are an abstraction (emulation) of physical hardware turning one server into many servers. The hypervisor allows multiple VMs to run on a single machine. Each VM includes a full copy of an operating system, one or more apps, necessary binaries and libraries - taking up tens of GBs. VMs can also be slow to boot.



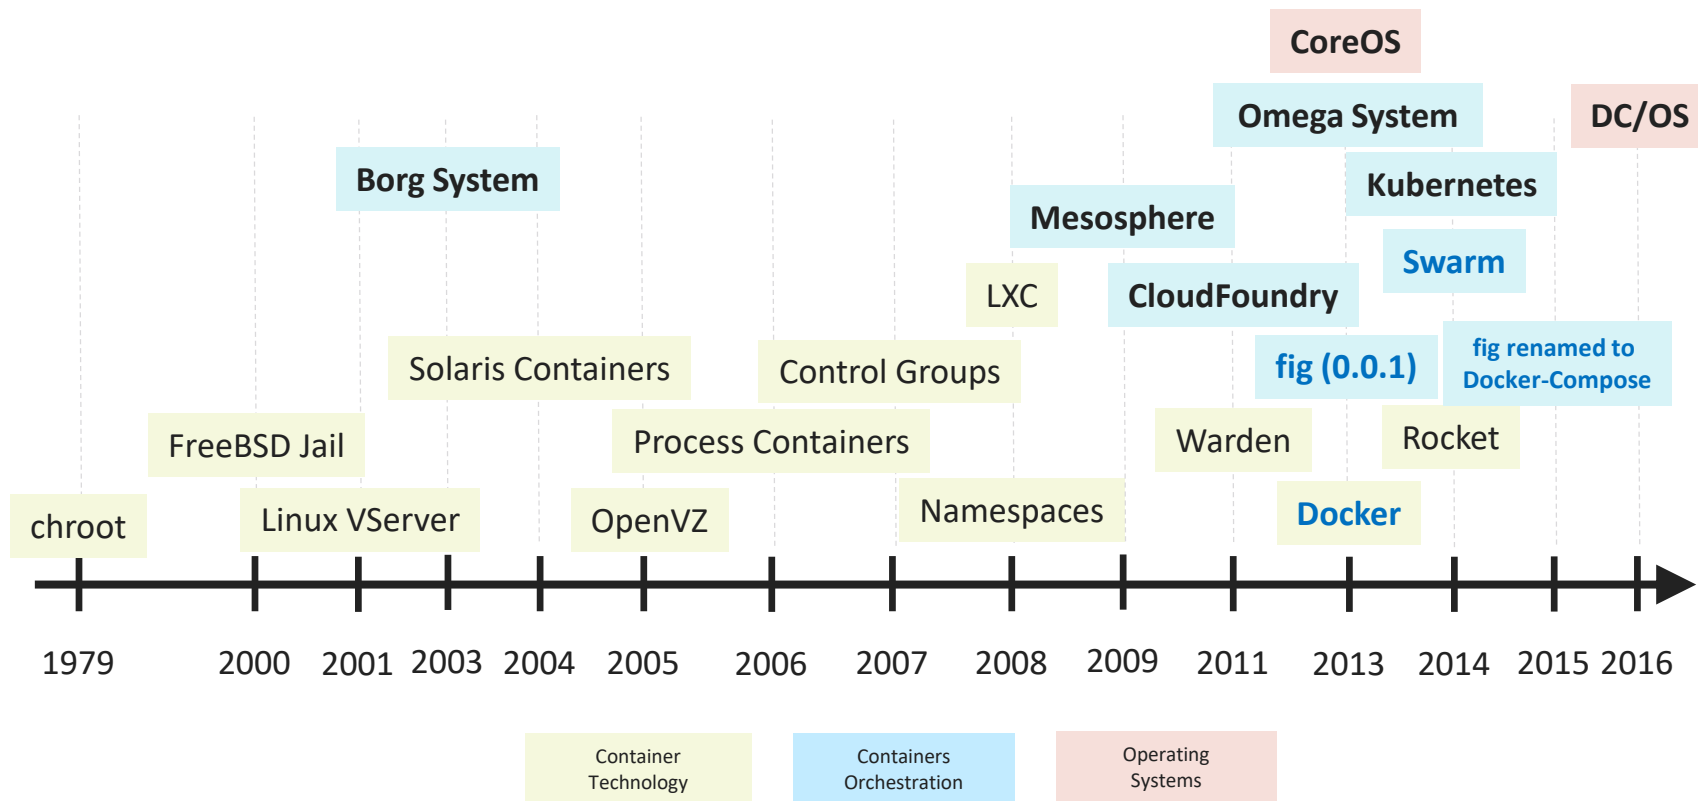
## CONTAINERS

Containers are an abstraction at the app layer that packages code and dependencies together (worker process). Multiple containers can run on the same machine and share the OS kernel with other containers, each running as isolated processes in user space. Containers take up less space than VMs (container images are typically tens of MBs in size), and start almost instantly.


# Virtual Machines vs Containers

	Container	Hypervisor
<b>Redundancy</b>	Only Application Code and Data Edited Files	Computation and Memory redundancy
<b>Flexibility</b>	Less Flexible, Linux can host linux based containers, Windows - windows based containers, MacOS uses Linux VM to run containers	More Flexible where Linux OS can host windows and vice verse
<b>Security</b>	Less Secure, DOS attacks can affect others Containers Container with root privileges could affect other Containers	Full Isolation, Hypervisor is responsible of limiting used resource by VM Cross Guest Access is prohibited
<b>Creating Time</b>	Almost instantaneous	Even Preexisting VMs need OS Load Time
<b>Performance</b>	Almost Native	Less than native due to middle ware
<b>Consolidation</b>	Limited by actual Application Usage	Limited By OS reserved
<b>Memory</b>	Allocation Memory	Allocated Files Disk, Allocated Files

# Containers History



# How OS Manages Containers



## Namespaces

When you run a container, Docker creates a set of *namespaces* for that container.

This provides a layer of isolation: each aspect of a container runs in its own namespace and does not have access outside of it.


Some of the namespaces that Docker Engine uses on Linux are:

- ✓ The **pid** namespace: Process isolation (PID: Process ID).
- ✓ The **net** namespace: Managing network interfaces (NET: Networking).
- ✓ The **ipc** namespace: Managing access to IPC resources (IPC: InterProcess Communication).
- ✓ The **mnt** namespace: Managing mount-points (MNT: Mount).
- ✓ The **uts** namespace: Isolating kernel and version identifiers. (UTS: Unix Timesharing System).

## [Manpages](#)



# How OS Manages Containers




## Control Groups

A key to running applications in isolation is to have them only use the resources you want. This ensures containers are good multi-tenant citizens on a host. Control groups allow Docker Engine to share available hardware resources to containers and, if required, set up limits and constraints. For example, limiting the memory available to a specific container.

- ✓ network.
- ✓ block i/o
- ✓ CPU
- ✓ Memory

# How OS Manages Containers



## chroot

**chroot** is simply isolation on the filesystem

## Union FS

operate by creating layers, making them very lightweight and fast. Docker Engine uses union file systems to provide the building blocks for containers. Docker Engine can make use of several union file system variants including: AUFS, btrfs, vfs, and DeviceMapper

# 12factor

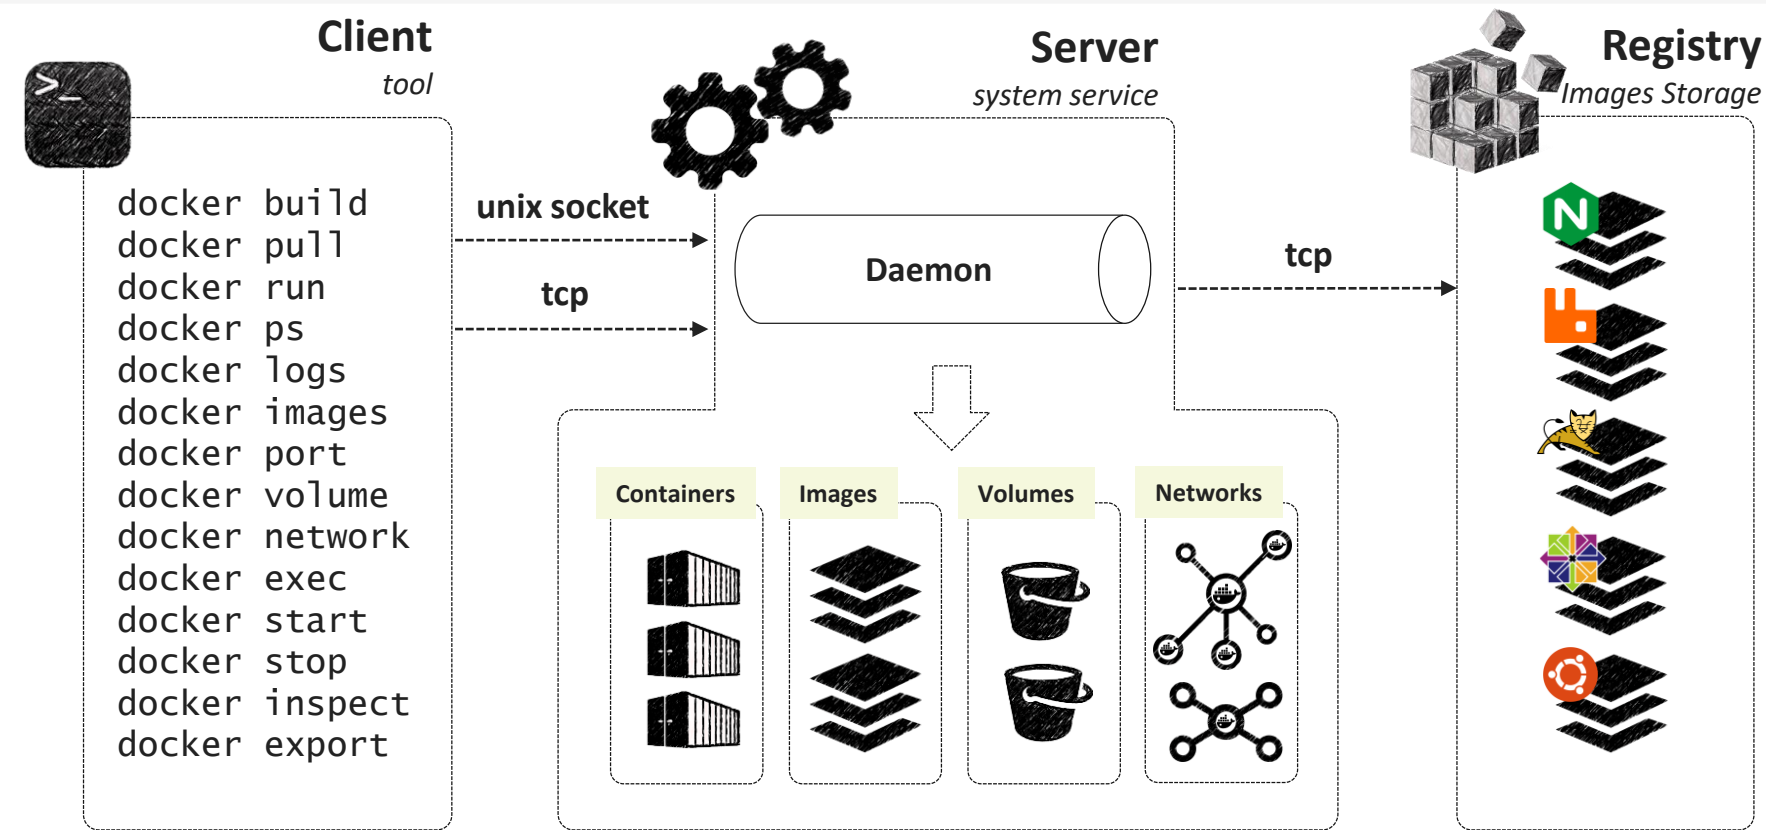
---

<https://github.com/docker/labs/tree/master/12factor>

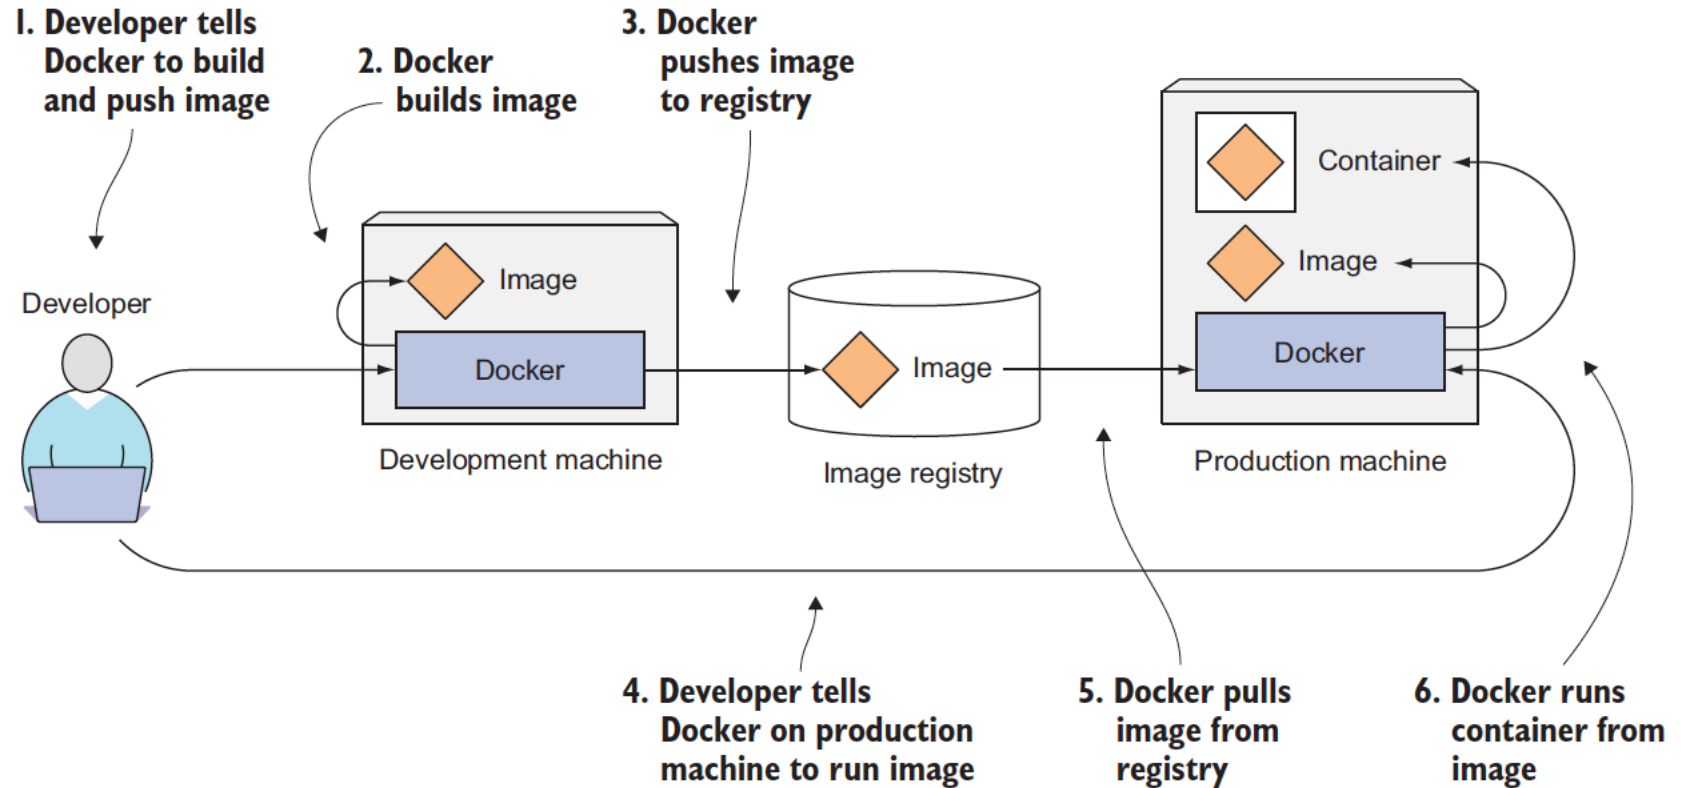
A great knowledge of what should be done to get cloud native application.

12 factor methodology is the result of their observations. As the name states, it presents 12 principles that will help application to be cloud ready, horizontally scalable, and portable.

# Docker Architecture



# Docker Workflow



# Docker Components

---

## The Docker daemon (dockerd)

listens for Docker API requests and manages Docker objects such as images, containers, networks, and volumes. A daemon can also communicate with other daemons to manage Docker services.

## The Docker client (docker)

is a tool to communicate with Docker daemon.

## Docker images

is a read-only template with instructions for creating a Docker container. Often, an image is *based on* another image, with some additional customization

## Docker containers

is a runnable instance of an image. You can create, start, stop, move, or delete a container using the Docker API or CLI

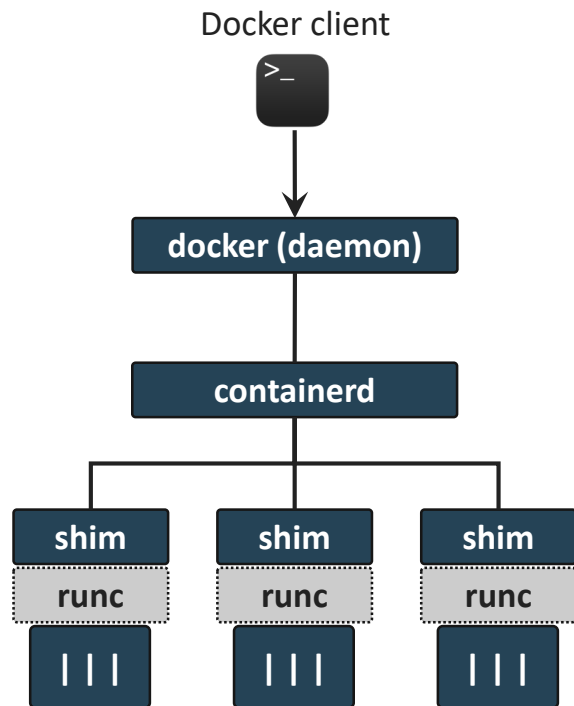
## Docker registries

stores Docker images

## Docker Services

allow you to scale containers across multiple Docker daemons, which all work together as a *swarm* with multiple *managers* and *workers*

# Docker Architecture



**(/usr/bin/)dockerd:** The Docker daemon itself. The highest level component in your list and also the only 'Docker' product listed. Provides all the nice UX features of Docker.

**(/usr/bin/)docker-containerd:** Also a daemon, listening on a Unix socket, exposes gRPC endpoints. Handles all the low-level container management tasks, storage, image distribution, network attachment, etc...

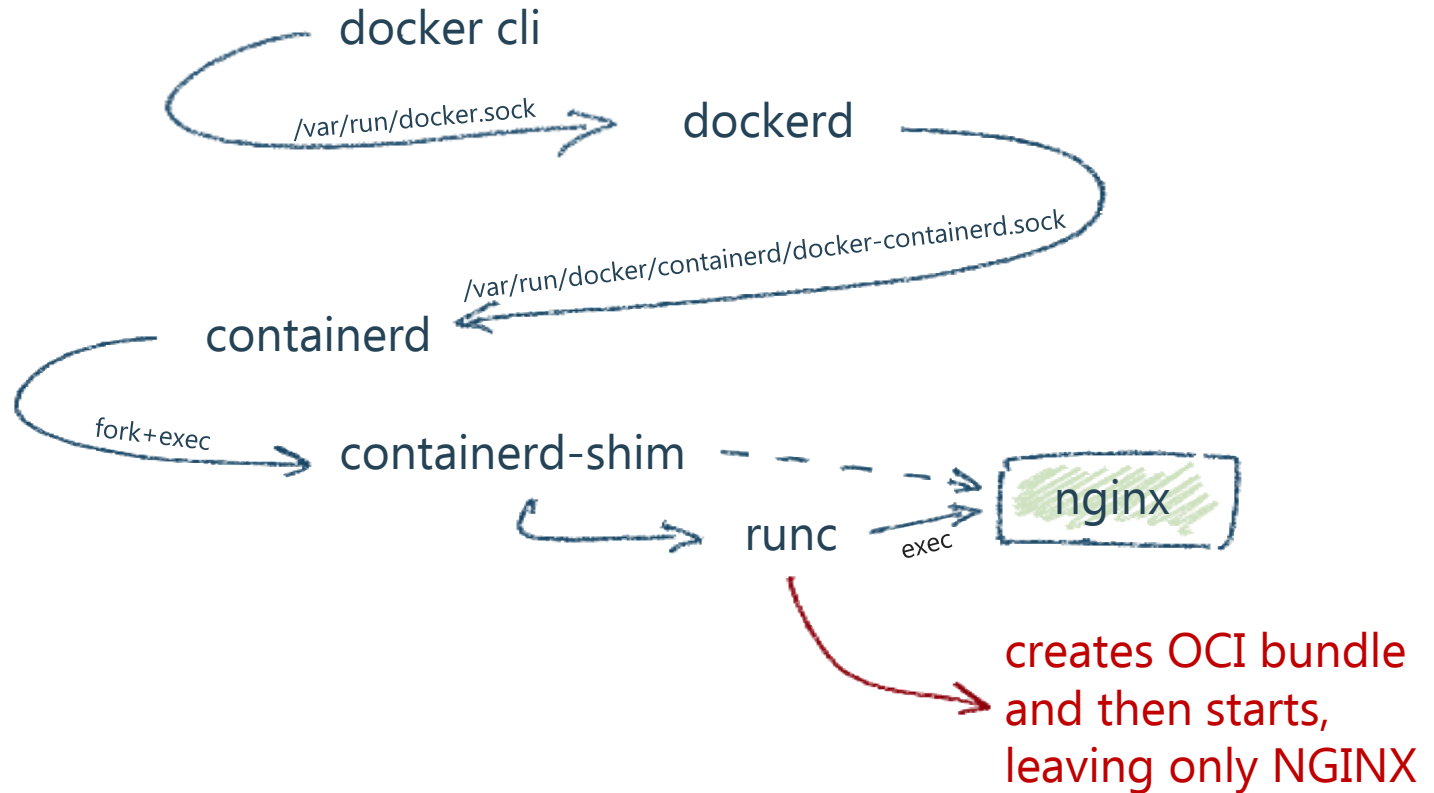
**(/usr/bin/)docker-containerd-ctr:** A lightweight CLI to directly communicate with containerd. Think of it as how 'docker' is to 'dockerd'.

**(/usr/bin/)docker-containerd-shim:** After runC actually runs the container, it exits (allowing us to not have any long-running processes responsible for our containers). The shim is the component which sits between containerd and runc to facilitate this.

**(/usr/bin/)docker-runc:** A lightweight binary for actually running containers. Deals with the low-level interfacing with Linux capabilities like cgroups, namespaces, etc...

**(/usr/bin/)docker-proxy:** A tool responsible for proxying container's ports to Host's interface

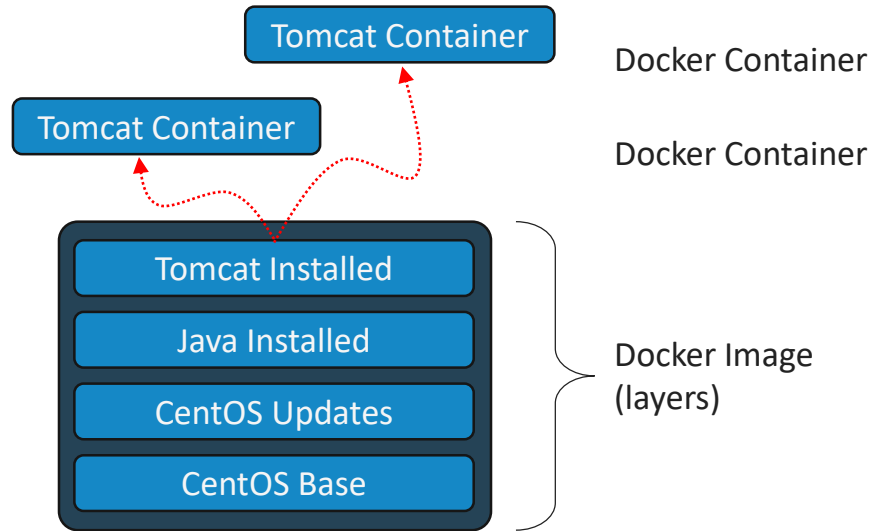
# How Docker Spawns Containers





# What is Image?

An image is an inert, immutable, file that's essentially a snapshot of a container. Images are created with the build command, and they'll produce a container when started with run. Images are stored in a Docker registry such as registry.hub.docker.com. Because they can become quite large, images are designed to be composed of layers of other images, allowing a minimal amount of data to be sent when transferring images over the network



## The copy-on-write (CoW) strategy

is a strategy of sharing and copying files for maximum efficiency. If a file or directory exists in a lower layer within the image, and another layer (including the writable layer) needs read access to it, it just uses the existing file. The first time another layer needs to modify the file (when building the image or running the container), the file is copied into that layer and modified. This minimizes I/O and the size of each of the subsequent layers. These advantages are explained in more depth below

# What is Image?

```
[vagrant@docker-host ~]$ docker pull sbeliakou/centos
```

```
Using default tag: latest
```

```
latest: Pulling from sbeliakou/centos
```

```
469cfcc7a4b3: Pull complete
```

```
b24939aa8741: Pull complete
```

```
[vagrant@docker-host ~]$ docker pull sbeliakou/ansible:2.6.1
```

```
2.6.1: Pulling from sbeliakou/ansible
```


```
469cfcc7a4b3: Already exists
```

```
f52300eb0803: Pull complete
```


```
22a01822f3c0: Pull complete
```

```
a3c99315580f: Pull complete
```

sbeliakou/training:centos-node




sbeliakou/centos:latest



sbeliakou/ansible:2.6.1



sbeliakou/ansible:2.6.2





# Docker installation and Configuration

- Installing Docker on CentOS
- How Docker Works

# Installing Docker Service on CentOS



```
Vagrant.configure("2") do |config|
  config.vm.box = "sbeliakou/centos"
  config.vm.box_version = "7.5"
  config.vm.network :private_network, ip: "192.168.56.15"

  config.vm.provision "shell", inline: <<-SHELL
    yum install -y yum-utils jq net-tools
    yum-config-manager --add-repo \
      https://download.docker.com/linux/centos/docker-ce.repo
    yum-config-manager --enable docker-ce-edge

    yum install -y docker-ce
    systemctl enable docker
    systemctl start docker
    usermod -aG docker vagrant
  SHELL
end
```



```
$ vagrant up
$ vagrant ssh
$ vagrant halt
$ vagrant destroy
```

# Installing Docker Service

---

## Add application user to docker group

The docker daemon binds to a Unix socket instead of a TCP port. By default that Unix socket is owned by the user root and other users can access it with sudo. For this reason, docker daemon always runs as the root user.

To avoid having to use sudo when you use the docker command, create a Unix group called docker and add users to it. When the docker daemon starts, it makes the ownership of the Unix socket read/writable by the docker group

```
[root@localhost ~]# usermod -aG docker vagrant
```

# Try Something Simple

```
[vagrant@localhost vagrant]$ docker run hello-world
```

```
Unable to find image 'hello-world:latest' locally  
latest: Pulling from library/hello-world  
9db2ca6ccae0: Pull complete  
Digest: sha256:4b8ff392a12ed9ea17784bd3c9a8b1fa3299cac44aca35a85c90c5e3c7afacdc  
Status: Downloaded newer image for hello-world:latest
```

Hello from Docker!

This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:

1. The Docker client contacted the Docker daemon.
2. The Docker daemon pulled the "hello-world" image from the Docker Hub.  
(amd64)
3. The Docker daemon created a new container from that image which runs the executable that produces the output you are currently reading.
4. The Docker daemon streamed that output to the Docker client, which sent it to your terminal.

To try something more ambitious, you can run an Ubuntu container with:

```
$ docker run -it ubuntu bash
```

Share images, automate workflows, and more with a free Docker ID:

<https://hub.docker.com/>

For more examples and ideas, visit:

<https://docs.docker.com/engine/userguide/>

# How It Works

Either by using the docker binary or via the API, the Docker client tells the Docker daemon to run a container.

```
[vagrant@localhost ~]$ docker run -i -t ubuntu bash
```

The Docker Engine client is launched using the **docker** tool with the **run** option running a new container.

The bare minimum the *docker client* needs to tell the Docker daemon to run the container is:

- ✓ What Docker image to build/create the container from, for example, `ubuntu`
- ✓ The command you want to run inside the container when it is launched, for example, `/bin/bash`

So what happens under the hood when we run this command?

In order, Docker Engine does the following:

- ✓ Pulls the **ubuntu** image.
- ✓ Creates a new container.
- ✓ Allocates a filesystem and mounts a read-write *layer*.
- ✓ Allocates a network / bridge interface.
- ✓ Sets up an IP address.
- ✓ Executes a process that you specify.
- ✓ Captures and provides application output.

You now have a running container! Now you can manage your container, interact with your application and then, when finished, stop and remove your container.

# Docker Daemon, Hub/Registry Commands

---

## Display Docker Version and Info

- `docker --version`
- `docker version`
- `docker version --format '{{.Server.Version}}'`
- `docker version --format '{{json .}}'`
- `docker info`
- `docker info --format '{{json .}}'`

## Working with Docker Hub/Registry:

- [docker login](#) - to login to a registry.
- [docker logout](#) - to logout from a registry.
- [docker search](#) - searches registry for image.
- [docker pull](#) - pulls an image from registry to local machine.
- [docker push](#) - pushes an image to the registry from local machine.



# Inspecting Docker Configuration

```
[vagrant@localhost ~]$ docker info
Containers: 0
  Running: 0
  Paused: 0
  Stopped: 0
Images: 0
Server Version: 18.06.0-ce
Storage Driver: overlay2
  Backing Filesystem: xfs
Logging Driver: json-file
Cgroup Driver: cgroupfs
Kernel Version: 3.10.0-862.9.1.el7.x86_64
Operating System: CentOS Linux 7 (Core)
Docker Root Dir: /var/lib/docker
Registry: https://index.docker.io/v1/
```



## Creating docker images

- **Build own Image with Dockerfile**
- **Tagging Images**
- **CMD/Entrypoint**
- **Build Arguments**
- **Multistage Build**

# Working with Images

---

## Lifecycle:

- [docker images](#) - shows all images.
- [docker import](#) - creates an image from a tarball.
- [docker build](#) - creates image from Dockerfile.
- [docker commit](#) - creates image from a container, pausing it temporarily if it is running.
- [docker rmi](#) - removes an image.
- [docker load](#) - loads an image from a tar archive as STDIN, including images and tags.
- [docker save](#) - saves an image to a tar archive stream to STDOUT with all parent layers, tags & versions.

## Info:

- [docker history](#) - shows history of image.
- [docker tag](#) - tags an image to a name (local or registry).

# Dockerfile



Dockerfile

```
FROM centos
LABEL maintainer="Siarhei Beliakou"
RUN yum install -y httpd web-assets-httpd && \
    yum clean all
RUN echo "my httpd container" > /var/www/html/index.html
# ADD/COPY index.html /var/www/html/
EXPOSE 80
ENTRYPOINT ["httpd"]
CMD ["-DFOREGROUND"]
```

```
# docker build [-t image_tag] [-f ./path/to/dockerfile] .
...
Status: Downloaded newer image for centos:latest
---> 49f7960eb7e4
...
Step 4/6 : RUN echo "my httpd container" > /var/www/html/index.html
---> Running in 37cc17740d0b
...
Successfully built 50a986f614d5
```

# Dockerfile Instructions

---

- [.dockerignore](#)
- [FROM](#) - Sets the Base Image for subsequent instructions.
- [RUN](#) - execute any commands in a new layer on top of the current image and commit the results.
- [CMD](#) - provide defaults for an executing container.
- [EXPOSE](#) - informs Docker that the container listens on the specified network ports at runtime.  
NOTE: does not actually make ports accessible.
- [ENV](#) - sets environment variable.
- [ADD](#) - copies new files, directories or remote file to container. Invalidates caches.
- [COPY](#) - copies new files or directories to container.
- [ENTRYPOINT](#) - configures a container that will run as an executable.
- [VOLUME](#) - creates a mount point for externally mounted volumes or other containers.
- [USER](#) - sets the user name for following RUN / CMD / ENTRYPOINT commands.
- [WORKDIR](#) - sets the working directory.
- [ARG](#) - defines a build-time variable.
- [ONBUILD](#) - adds a trigger instruction when the image is used as the base for another build.
- [STOPSIGNAL](#) - sets the system call signal that will be sent to the container to exit.
- [LABEL](#) - apply key/value metadata to your images, containers, or daemons.

# Tagging Images

---

```
$ docker build -t myhttpd .  
...  
Successfully built 50a986f614d5  
Successfully tagged myhttpd:latest
```

```
$ docker tag 50a986f614d5 sbeliakou/myhttpd
```

```
$ docker tag myhttpd sbeliakou/myhttpd:1.0
```

```
$ docker tag myhttpd:latest sbeliakou/myhttpd:latest
```

# CMD/ENTRYPOINT



```
FROM busybox
ENTRYPOINT ["echo", "hello"]
```

```
# docker build -t cmd_entrypoint:0.1 .
```

```
# docker run cmd_entrypoint:0.1
hello
```

```
# docker run cmd_entrypoint:0.1 world
hello world
```



```
FROM busybox
ENTRYPOINT echo hello
```

```
# docker build -t cmd_entrypoint:0.2 .
```

```
# docker run cmd_entrypoint:0.2
hello
```

```
# docker run cmd_entrypoint:0.2 world
hello
```



```
FROM busybox
CMD ["echo", "hello"]
```

```
# docker build -t cmd_entrypoint:0.3 .
```

```
# docker run cmd_entrypoint:0.3
hello
```

```
# docker run cmd_entrypoint:0.3 world
docker: Error response from daemon: OCI runtime create
failed: container_linux.go:348: starting container
process caused "exec: \"world\": executable file not
found in $PATH": unknown.
```



```
FROM busybox
CMD echo hello
```

```
# docker build -t cmd_entrypoint:0.4 .
```

```
# docker run cmd_entrypoint:0.4
hello
```

```
# docker run cmd_entrypoint:0.4 world
docker: Error response from daemon: OCI runtime create
failed: container_linux.go:348: starting container
process caused "exec: \"world\": executable file not
found in $PATH": unknown.
```

# CMD/ENTRYPOINT



```
FROM busybox
ENTRYPOINT echo
CMD hello
```

```
# docker build -t cmd_entrypoint:0.5 .
```

```
# docker run cmd_entrypoint:0.5
```

```
# docker run cmd_entrypoint:0.5 world
```



```
FROM busybox
ENTRYPOINT echo
CMD ["hello"]
```

```
# docker build -t cmd_entrypoint:0.7 .
```

```
# docker run cmd_entrypoint:0.7
```

```
# docker run cmd_entrypoint:0.7 world
```



```
FROM busybox
ENTRYPOINT ["echo"]
CMD hello
```

```
# docker build -t cmd_entrypoint:0.6 .
```

```
# docker run cmd_entrypoint:0.6
/bin/sh -c hello
```

```
# docker run cmd_entrypoint:0.6 world
world
```



```
FROM busybox
ENTRYPOINT ["echo"]
CMD ["hello"]
```

```
# docker build -t cmd_entrypoint:0.8 .
```

```
# docker run cmd_entrypoint:0.8
hello
```

```
# docker run cmd_entrypoint:0.8 world
world
```



# CMD/ENTRYPOINT. The Best Choice



```
...  
ENTRYPOINT ["/command/to/be/executed"]  
CMD ["default", "args"]
```



```
docker run image arg1 arg2 ...
```



```
entrypoint arg1 arg2 ...
```

```
docker run image
```



```
entrypoint cmd
```

# Dockerfile Example: Custom Packer Image



```
FROM centos:7
ENV PACKER_VERSION 1.2.3
RUN yum install -y \
    epel-release \
    yum-plugin-ovl \
    wget unzip \
    rsync \
    openssh openssh-clients && \
    yum install -y python-pip && \
    yum clean all

RUN wget -q https://releases.hashicorp.com/...${PACKER_VERSION}_linux_amd64.zip && \
    unzip -q packer_${PACKER_VERSION}_linux_amd64.zip -d /bin/ && \
    rm -f packer_${PACKER_VERSION}_linux_amd64.zip
RUN pip install -U ansible ansible-modules-hashivault

RUN useradd packer

USER packer
ENV USER packer

ENTRYPOINT ["/bin/packer"]
CMD ["--version"]
```

## Some More Examples



```
FROM java:8-jre
```

```
ADD customer-contact-service.jar /  
EXPOSE 4040
```

```
CMD ["java", "-jar", "/customer-contact-service.jar"]
```



```
FROM java:8
```

```
ENV PORT=8080  
EXPOSE 8080  
COPY wiremock /wiremock
```

```
CMD ["/wiremock/bin/startServer.sh"]
```

# Docker Base Images

- **scratch** – this is the ultimate base image and it has 0 files and 0 size.
- **busybox** – a minimal Unix weighing in at 2.5 MB and around 10000 files.
- **debian:jessie** – the latest Debian is 122 MB and around 18000 files.
- **alpine:latest** – Alpine Linux, only 8 MB in size and has access to a package repository



```
FROM scratch
```

```
ADD centos.tar.gz /
```

```
RUN yum install -y epel-release && \  
    yum update -y && \  
    yum clean all
```

```
LABEL architecture="amd64" \  
      OS="CentOS" \  
      License=GPLv2 \  
      maintainer="Siarhei Beliakou (sbeliakou@gmail.com)"
```

```
## Default command
```

```
CMD ["/bin/bash"]
```

[More details](#)

# Useful Links and Examples

---

**Many Samples:** [docs.docker.com/samples/](https://docs.docker.com/samples/)

**Examples:** [#dockerfile-examples](https://github.com/dockerfile-examples)

**Jenkins:** [Dockerfile-alpine](https://github.com/dockerfile-alpine)

**Tomcat:** [jre8-alpine/Dockerfile](https://github.com/jre8-alpine/Dockerfile)

**OpenJDK:** [jdk/alpine/Dockerfile](https://github.com/jdk/alpine/Dockerfile)

**AmazonLinux:** [2018.03/Dockerfile](https://github.com/2018.03/Dockerfile)

**Another example is here:**

[#define-a-container-with-dockerfile](https://github.com/define-a-container-with-dockerfile)

**Dockerfile Reference:**

[#/dockerfile-reference](https://github.com/dockerfile-reference)

**Best practices for writing Dockerfiles:**

[docs.docker.com/dockerfile best-practices/](https://docs.docker.com/dockerfile_best-practices/)

**Docker CLI Reference:**

[commandline/cli/](https://docs.docker.com/engine/reference/commandline/cli/)

# Build With Arguments



```
ARG BASE_IMAGE
FROM ${BASE_IMAGE}
...
```

```
# docker build --build-arg BASE_IMAGE=hashicorp/terraform:0.10.8.
...
```

```
# docker build --build-arg BASE_IMAGE=hashicorp/terraform:0.11.7.
...
```



```
...
ARG BUILD_ID
LABEL build_id="${BUILD_ID}"
...
```

# Multistage Build



```
FROM maven:3.3-jdk-8 as builder
COPY . /build/
WORKDIR /build
RUN mvn clean install
```

```
FROM openjdk:8-jre
COPY --from=builder /build/target/demoapp.jar /opt/
EXPOSE 8080
ENTRYPOINT ["java", "-jar", "/opt/demoapp.jar"]
```

```
# docker build -t myapp .
```

```
...
```

```
# docker build -t myapp --target builder .
```

```
...
```

Explore example here: [sbeliakou/springboot example](https://sbeliakou.github.io/springboot-example/)



## RUNNING CONTAINERS

- Running in detached mode
- Exposing Ports
- Managing Restart Policy
- Changing Workspace Directory
- Changing Runtime User
- Providing Environment Variables



# Working with Containers

---

## Lifecycle:

- [docker create](#) - creates a container but does not start it.
- [docker rename](#) - allows the container to be renamed.
- [docker run](#) - creates and starts a container in one operation.
- [docker rm](#) - deletes a container.
- [docker update](#) - updates a container's resource limits.

## Starting and Stopping:

- [docker start](#) - starts a container so it is running.
- [docker stop](#) - stops a running container.
- [docker restart](#) - stops and starts a container.
- [docker pause](#) - pauses a running container, "freezing" it in place.
- [docker unpause](#) - will unpause a running container.
- [docker wait](#) - blocks until running container stops.
- [docker kill](#) - sends a SIGKILL to a running container.
- [docker attach](#) - will connect to a running container.

# Working with Containers

---

## Info:

- [docker ps](#) - shows running containers.
- [docker logs](#) - gets logs from container.
- [docker inspect](#) - looks at all the info on a container.
- [docker events](#) - gets events from container.
- [docker port](#) - shows public facing port of container.
- [docker top](#) - shows running processes in container.
- [docker stats](#) - shows containers' resource usage statistics.
- [docker diff](#) - shows changed files in the container's FS.

## Import / Export:

- [docker cp](#) - copies files or folders between a container and the local filesystem.
- [docker export](#) - turns container filesystem into tarball archive stream to STDOUT.

## Executing Commands:

- [docker exec](#) - to execute a command in container.

# Running the Container

```
# docker run 50a986f614d5
```

```
^C
```

```
# docker run -d myhttpd:1.0
```

```
9f761335efe268e9a82c4828d8f4be67b5824eb3266e8ba311343a7da45c67ff
```

```
# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
9f761335efe2	50a986f614d5	"/bin/sh -c 'httpd -..."	5 seconds ago	Up 4 seconds	80/tcp	trusting_kilby

```
# docker run -P -d myhttpd:1.0
```

```
74954ff14ec5e53ac9925bfd2873c654fe8978657764b4162ac494fc9afaab9f
```

```
# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
74954ff14ec5	myhttpd:1.0	"/bin/sh -c 'httpd -..."	9 seconds ago	Up 18 seconds	0.0.0.0:32768->80/tcp	fervent_noyce
9f761335efe2	myhttpd:1.0	"/bin/sh -c 'httpd -..."	3 minutes ago	Up 3 minutes	80/tcp	trusting_kilby

```
# curl localhost:32768
```

```
my httpd container
```

```
# docker run -d -p 8081:80 --name h8081 myhttpd:1.0
```

```
fca7f4525bc618e7c503b73bfa680c055300e8b5c767d48e33669831e0bc5bec
```

```
# docker run -d -p 127.0.0.1:8082:80 --name h8082 myhttpd:1.0
```

```
014e5efa5ca90d9b7e50eebef3c7f020f08f0b5238f98420681ee348a4097829
```

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.Ports}}" -n2
```

NAMES	IMAGE	CONTAINER ID	PORTS
h8081	myhttpd:1.0	fca7f4525bc6	0.0.0.0:8081->80/tcp
h8082	myhttpd:1.0	014e5efa5ca9	127.0.0.1:8082->80/tcp

formatting

# Running the Container: Restarting Policy

```
# docker run -d --restart=always --name sleeper centos sleep 5
6c3d24b3f89f13de92e710fcbb5b343b4cb81e7454ecc79445e94f0c5ba31a49
```

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1
```

NAMES	IMAGE	CONTAINER ID	CREATED	STATUS
sleeper	centos	33675bfff7f47	2 seconds ago	Up 1 second

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1
```

NAMES	IMAGE	CONTAINER ID	CREATED	STATUS
sleeper	centos	33675bfff7f47	11 seconds ago	Up 4 seconds

```
# docker ps --format "table {{.Names}}\t{{.Image}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1
```

NAMES	IMAGE	CONTAINER ID	CREATED	STATUS
sleeper	centos	33675bfff7f47	About a minute ago	Restarting (0) 6 seconds ago

```
# docker inspect -f "{{ .RestartCount }}" sleeper
19
```

[More details](#)

Policy	Result
no	Do not automatically restart the container when it exits. This is the default.
on-failure[:max-retries]	Restart only if the container exits with a non-zero exit status. Optionally, limit the number of restart retries the Docker daemon attempts.
always	Always restart the container regardless of the exit status. When you specify always, the Docker daemon will try to restart the container indefinitely. The container will also always start on daemon startup, regardless of the current state of the container.
unless-stopped	Always restart the container regardless of the exit status, including on daemon startup, except if the container was put into a stopped state before the Docker daemon was stopped.

# Running Containers in Interactive Mode

```
[root@localhost ~]# docker run centos cat /etc/redhat-release
CentOS Linux release 7.5.1804 (Core)
```

```
[root@localhost ~]# docker run ubuntu cat /etc/lsb-release
DISTRIB_ID=Ubuntu
DISTRIB_RELEASE=18.04
DISTRIB_CODENAME=bionic
DISTRIB_DESCRIPTION="Ubuntu 18.04.1 LTS"
```

```
[root@localhost ~]# docker run -it centos bash
[root@dfc1b0d4f6a5 /]# cat /etc/redhat-release
CentOS Linux release 7.5.1804 (Core)
[root@dfc1b0d4f6a5 /]# yum install epel-release
...
[root@dfc1b0d4f6a5 /]# exit
exit
[root@localhost ~]#
```

```
[root@localhost ~]# docker ps --format "table {{.ID}}\t{{.RunningFor}}\t{{.Status}}" -n1 -a
CONTAINER ID      CREATED          STATUS
dfc1b0d4f6a5      7 minutes ago   Exited (1) 1 minutes ago
```

```
[root@localhost ~]# docker rm dfc1b0d4f6a5
dfc1b0d4f6a5
```

```
[root@localhost ~]# docker run --rm centos cat /etc/redhat-release
CentOS Linux release 7.5.1804 (Core)
```

# Executing Commands Inside Running Container

```
# docker run -d centos sleep infinity
9626a94669c935ea140bcb9ea83339bd325b28195fc088a690b620eb12902b33
```

```
# docker ps -l
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
9626a94669c9	centos	"sleep infinity"	11 seconds ago	Up 10 seconds		priceless_einstein

```
# docker exec -it 9626a94669c9 bash
```

```
[root@9626a94669c9 /]# ps -ef
```

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	1	0	0	13:43	?	00:00:00	sleep infinity
root	34	0	2	13:48	pts/0	00:00:00	bash
root	47	34	0	13:48	pts/0	00:00:00	ps -ef

```
[root@9626a94669c9 /]# exit
exit
#
```

# Stopping/Deleting Containers

```
# docker ps --format "table {{.Image}}\t{{.Names}}\t{{.ID}}\t{{.RunningFor}}\t{{.Status}}"
```

IMAGE	NAMES	CONTAINER ID	CREATED	STATUS
myhttpd:1.0	h8081	fca7f4525bc6	About an hour ago	Up About an hour
myhttpd:1.0	h8082	014e5efa5ca9	About an hour ago	Up About an hour

```
# docker stop h8082  
014e5efa5ca9
```

```
# docker rm 014e5efa5ca9  
014e5efa5ca9
```

```
# docker rm $(docker stop 014e5efa5ca9)  
014e5efa5ca9
```

```
# docker rm $(docker stop $(docker ps -a -q))  
014e5efa5ca9  
fca7f4525bc6
```

☆ 

```
# docker container prune  
WARNING! This will remove all stopped containers.  
Are you sure you want to continue? [y/N]
```

☆ 

```
# docker image prune  
WARNING! This will remove all dangling images.  
Are you sure you want to continue? [y/N]  
you want to continue? [y/N]
```

# Changing Container's Build Defaults

## 1. Default User

```
# docker run -it jenkins id  
uid=1000(jenkins) gid=1000(jenkins) groups=1000(jenkins)
```

```
# docker run -it --user 0 jenkins id  
uid=0(root) gid=0(root) groups=0(root)
```

**--user == -u**

```
# docker run -it --user 1000:0 jenkins id  
uid=1000(jenkins) gid=0(root) groups=0(root)
```

```
# docker run -it --group-add 123 jenkins id  
uid=1000(jenkins) gid=1000(jenkins) groups=1000(jenkins),123
```

## 2. Default Workdir

```
# docker run -it jenkins pwd  
/
```

```
# docker run -it --workdir /var/jenkins_home jenkins pwd  
/var/jenkins_home
```

**--workdir == -w**

```
# docker run -it -w $(pwd) -v $(pwd):$(pwd) maven clean package
```



# Changing Container's Build Defaults

## 3. Default Entrypoint

```
# docker run -it myhttpd:1.0 bash
# echo $?
1
```

```
# docker run -it --entrypoint=/bin/bash myhttpd:1.0
[root@e4e7976cf838 /]# ps -ef
UID          PID    PPID  C STIME TTY          TIME CMD
root           1         0  1 19:27 pts/0        00:00:00 /bin/bash
root          14         1  0 19:27 pts/0        00:00:00 ps -ef
[root@e4e7976cf838 /]# exit
#
```

## 4. Environment Variables

```
# docker run -it -e MYVAR="My Variable" centos env | grep MYVAR
MYVAR=My Variable
```

```
# docker run -it --env-file <(env | grep ARM | cut -f1 -d=) centos env | grep ARM
ARM_SUBSCRIPTION_ID=64a3f30f-xxx-xxxx-xxxx-yyyyfe63fe9a
ARM_TENANT_ID=bd5c6713-xxx-yyyy-xxxx-78f2d078e543
ARM_CLIENT_SECRET=xxxxxxxxxxxxxxxxxxxx
ARM_CLIENT_ID=808f38ed-xxxx-xxxx-yyyy-ebbce91bcfee
```



## MOUNTS

- Mounting Data from Host

# Mounting Data from Host

We have a directory on the host with valuable data which we need to share into the container

```
# ls -l ./
total 40
-rw-r--r--  1 sbeliakou  wheel  119 July 25 22:05 index.html

# docker run -d -P -v $(pwd):/var/www/html httpd
# docker run -d -P -v $(pwd):/var/www/html:ro httpd
# docker run -d -p 127.0.0.1:8080:80 -v $(pwd):/var/www/html:ro httpd
```

The -v flag can also be used to mount a single file - instead of *just* directories - from the host machine

```
# docker run --rm -it -v ~/.bash_history:/root/.bash_history ubuntu /bin/bash
```

```
# ls -l ./
total 40
-rw-r--r--  1 root  root  119 July 25 22:05 index.html

# docker run -P -v $(pwd):$(pwd) -w $(pwd) busybox ls -l
total 40
-rw-r--r--  1 root  root  119 July 25 22:05 index.html
```

# Mounting Data from Host

```
# ls
playbook.yml

# which ansible-playbook
/usr/bin/which: no ansible-playbook in (/sbin:/bin:/usr/sbin:/usr/bin)

# alias ansible-playbook='docker run -v $(pwd):$(pwd) -w $(pwd) ansible:2.6.1'

# ansible-playbook playbook.yml -vv
PLAYBOOK: playbook.yml *****
1 plays in playbook.yml

PLAY [localhost] *****

TASK [Gathering Facts] *****
task path: /vagrant/playbook.yml:1
ok: [localhost]

TASK [debug] *****
task path: /vagrant/playbook.yml:4
ok: [localhost] => {
    "ansible_host": "localhost"
}

PLAY RECAP *****
localhost                : ok=2    changed=0    unreachable=0    failed=0
```



## GETTING LOGS FROM CONTAINERS

- In-container Logs
- Container Log Drivers

# Container Logs

By default, **docker logs** shows the command's **STDOUT** and **STDERR**

```
# docker run hello-world
# docker logs $(docker ps -aq1)
```

Let's check container from *myhttpd:latest* image:

```
# docker run -d -P myhttpd:latest
791c65b5aed05a11a40a953779675b5b94d090e691730c5c0bceaa33143fcbc4

# docker logs $(docker ps -lq)
AH00558: httpd: Could not reliably determine the server's fully qualified domain name,
using 172.17.0.17. Set the 'ServerName' directive globally to suppress this message
```

```
# curl localhost:$(docker port $(docker ps -lq) | cut -d: -f2)
my httpd container

# docker logs $(docker ps -lq)
AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using
172.17.0.17. Set the 'ServerName' directive globally to suppress this message
```

# Container Logs

Let's create "Dockerfile" file with following content



```
FROM centos
LABEL maintainer="Siarhei Beliakou"
RUN yum install -y httpd web-assets-httpd && yum clean all
RUN echo "logs are sending to stdout" > /var/www/html/index.html
# ADD index.html /var/www/html/
RUN ln -s /dev/stdout /var/log/httpd/access_log && \
    ln -s /dev/stderr /var/log/httpd/error_log
EXPOSE 80
CMD httpd -DFOREGROUND
```



```
FROM myhttpd:1.0
RUN echo "logs are sending to stdout" > /var/www/html/index.html
RUN ln -s /dev/stdout /var/log/httpd/access_log && \
    ln -s /dev/stderr /var/log/httpd/error_log
```

And build it:

```
# docker build -t myhttpd:2.0 .
...
Successfully built 5ffd0ffc1780
Successfully tagged myhttpd:2.0
```

# Container Logs

```
# docker run -d -P myhttpd:2.0
```

```
1b0c3a82d0687519ffcd2ee06d345a4d3ad8d475dc0d7710fb279bdd836e5e88
```

```
# docker logs $(docker ps -lq)
```

```
AH00558: httpd: Could not reliably determine the server's fully qualified domain name,  
using 172.17.0.17. Set the 'ServerName' directive globally to suppress this message
```

```
# curl localhost:$(docker port $(docker ps -lq) | cut -d: -f2)
```

```
logs are sending to stdout
```

```
# docker logs $(docker ps -lq)
```

```
AH00558: httpd: Could not reliably determine the server's fully qualified domain name, using  
172.17.0.18. Set the 'ServerName' directive globally to suppress this message
```

```
172.17.0.1 - - [29/Jul/2018:22:07:24 +0000] "GET / HTTP/1.1" 200 19 "-" "curl/7.29.0"
```



# Configuring Log Driver

```
# docker run -d -P --name=myhttpd --log-driver=journald myhttpd:2.0  
5748f3078b647c43a335bdc1f8bc7c8d9db31a491e635effd58b31b1429c8ee7
```

```
# curl localhost:$(docker port $(docker ps -lq) | cut -d: -f2)  
logs are sending to stdout
```

```
# journalctl -b CONTAINER_NAME=myhttpd  
-- Logs begin at Sat 2018-07-28 19:14:07 BST, end at Sun 2018-07-29 23:19:52 BST. --  
Jul 29 23:19:29 localhost.localdomain 5748f3078b64[2806]: AH00558: httpd: Could not reliably  
determine the server's fully qualified domain name, using 172.17.  
Jul 29 23:19:50 localhost.localdomain 5748f3078b64[2806]: 172.17.0.1 - - [29/Jul/2018:22:19:50  
+0000] "GET / HTTP/1.1" 200 19 "-" "curl/7.29.0"
```

# Configuring Log Driver

```
# docker run -d -P --log-driver=journald --log-opt tag=myhttpd myhttpd:latest  
0555d7a41ab6af098dfea296e2fa7b4f1c6b73424e2156c387930b13bfcefb24
```

```
# curl localhost:$(docker port $(docker ps -lq) | cut -d: -f2)  
logs are sending to stdout
```

```
# journalctl -b CONTAINER_TAG=myhttpd  
-- Logs begin at Sat 2018-07-28 19:14:07 BST, end at Sun 2018-07-29 23:27:27 BST. --  
Jul 29 23:26:02 localhost.localdomain myhttpd[2806]: AH00558: httpd: Could not reliably determine  
the server's fully qualified domain name, using 172.17.0.21.  
Jul 29 23:26:26 localhost.localdomain myhttpd[2806]: 172.17.0.1 - - [29/Jul/2018:22:26:26 +0000]  
"GET / HTTP/1.1" 200 19 "-" "curl/7.29.0"
```

# Supported Log Drivers

Driver	Description
<a href="#">none</a>	No logs are available for the container and docker logs does not return any output.
<a href="#">json-file</a>	The logs are formatted as JSON. The default logging driver for Docker.
<a href="#">syslog</a>	Writes logging messages to the syslog facility. The syslog daemon must be running on the host machine.
<a href="#">journald</a>	Writes log messages to journald. The journald daemon must be running on the host machine.
<a href="#">gelf</a>	Writes log messages to a Graylog Extended Log Format (GELF) endpoint such as Graylog or Logstash.
<a href="#">fluentd</a>	Writes log messages to fluentd (forward input). The fluentd daemon must be running on the host machine.
<a href="#">awslogs</a>	Writes log messages to Amazon CloudWatch Logs.
<a href="#">splunk</a>	Writes log messages to splunk using the HTTP Event Collector.
<a href="#">etwlogs</a>	Writes log messages as Event Tracing for Windows (ETW) events. Only available on Windows platforms.
<a href="#">gcplogs</a>	Writes log messages to Google Cloud Platform (GCP) Logging.
<a href="#">logentries</a>	Writes log messages to Rapid7 Logentries.

# Dockerd: Configuring Log Driver

```
# docker info --format '{{.LoggingDriver}}'  
json-file  
  
# cat << EOF > /etc/docker/daemon.json  
{  
  "log-driver": "journald"  
}  
EOF  
  
# systemctl daemon-reload  
# systemctl restart docker.service  
# docker info --format '{{.LoggingDriver}}'  
journald
```

```
# docker run -d -P --name=myhttpd_2.0 myhttpd:2.0
```

```
# journalctl -b CONTAINER_NAME=myhttpd_2.0
```

# Demo

---

<https://github.com/dbuhtiyarov/docker-demo/tree/master/java/DemoSpringBootApplication>



# Thank you for your attention!

Dmitry Buhtiyarov, Siarhei Beliakou  
2018

